

Computational Intractability (II) & Pseudorandomness (I)

601.642/442: Modern Cryptography

Fall 2022

Asymptotic Cost of an Attack

- As we saw in the last class, exponential-time (e.g., $\Theta(2^n)$) brute-force attacks do not scale well as we increase the length of the key or the security parameter. We view them as *infeasible*.
- To rule out only *feasible* attacks, we will focus on polynomial-time attacks (e.g., n^2, n^5 , etc) that scale reasonably well (especially if exponent is small). Recall that polynomial-time algorithms are also referred to as *efficient* algorithms.
- **Our goal:** Ensure that no polynomial-time attack can successfully break security.
- A useful property of polynomial-time algorithms is **closure**: *repeating a poly-time algorithm polynomial times is still polynomial time!*

Efficient Adversary

- **Computational Security:** Security against adversaries that run in polynomial-time.

Efficient Adversary

- **Computational Security:** Security against adversaries that run in polynomial-time.
- The adversary might be *randomized*.

Efficient Adversary

- **Computational Security:** Security against adversaries that run in polynomial-time.
- The adversary might be *randomized*.
- It might also use a *different* algorithm for each input size, each of which might be efficient. This still counts as efficient since he is using polynomial-time resources!

Efficient Adversary

- **Computational Security:** Security against adversaries that run in polynomial-time.
- The adversary might be *randomized*.
- It might also use a *different* algorithm for each input size, each of which might be efficient. This still counts as efficient since he is using polynomial-time resources!
- We call this a **non-uniform** adversary since the algorithm is not uniform across all input sizes.

Non-Uniform PPT

Definition (Non-Uniform PPT)

A non-uniform probabilistic polynomial time Turing machine A is a sequence of probabilistic machines $A = \{A_1, A_2, \dots\}$ for which there exists a polynomial $p(\cdot)$ such that for every $A_i \in A$, the description size $|A_i|$ and the running time of A_i are at most $p(i)$. We write $A(x)$ to denote the distribution obtained by running $A_{|x|}(x)$.

- Our adversary will usually be a non-uniform PPT Turing machine.
(most general)

Success Probability of an Attack

- Since our adversary is **probabilistic**, it is not enough to consider only the running time of an attack.

Success Probability of an Attack

- Since our adversary is **probabilistic**, it is not enough to consider only the running time of an attack.
- For example, consider an attacker who just tries to guess a victim's secret key, making a single guess. This attack is extremely cheap, but it still has a nonzero chance (e.g., 2^{-128}) of breaking security.

Success Probability of an Attack

- Since our adversary is **probabilistic**, it is not enough to consider only the running time of an attack.
- For example, consider an attacker who just tries to guess a victim's secret key, making a single guess. This attack is extremely cheap, but it still has a nonzero chance (e.g., 2^{-128}) of breaking security.
- Essentially, for security, we don't need to worry about the following:

Success Probability of an Attack

- Since our adversary is **probabilistic**, it is not enough to consider only the running time of an attack.
- For example, consider an attacker who just tries to guess a victim's secret key, making a single guess. This attack is extremely cheap, but it still has a nonzero chance (e.g., 2^{-128}) of breaking security.
- Essentially, for security, we don't need to worry about the following:
 - Attacks that are as expensive as a brute-force attack.

Success Probability of an Attack

- Since our adversary is **probabilistic**, it is not enough to consider only the running time of an attack.
- For example, consider an attacker who just tries to guess a victim's secret key, making a single guess. This attack is extremely cheap, but it still has a nonzero chance (e.g., 2^{-128}) of breaking security.
- Essentially, for security, we don't need to worry about the following:
 - Attacks that are as expensive as a brute-force attack.
 - Attacks whose success probability is as low as a blind-guess attack.

Success Probability of an Attack

- Since our adversary is **probabilistic**, it is not enough to consider only the running time of an attack.
- For example, consider an attacker who just tries to guess a victim's secret key, making a single guess. This attack is extremely cheap, but it still has a nonzero chance (e.g., 2^{-128}) of breaking security.
- Essentially, for security, we don't need to worry about the following:
 - Attacks that are as expensive as a brute-force attack.
 - Attacks whose success probability is as low as a blind-guess attack.
- While an attack with success probability 2^{-128} should not really count as an attack, one with success probability $1/2$ should. Where should we draw the line?

Success Probability of an Attack

Probability	Equivalent
2^{-10}	full house in 5-card poker
2^{-20}	royal flush in 5-card poker
2^{-28}	you win this week's Powerball jackpot
2^{-40}	royal flush in 2 consecutive poker games
2^{-60}	the next meteorite that hits Earth lands on this slide

Success Probability of an Attack

Probability	Equivalent
2^{-10}	full house in 5-card poker
2^{-20}	royal flush in 5-card poker
2^{-28}	you win this week's Powerball jackpot
2^{-40}	royal flush in 2 consecutive poker games
2^{-60}	the next meteorite that hits Earth lands on this slide

- While these examples are good to get some intuition, we need more concrete way to draw the line between *reasonable* and *unreasonable*.

Success Probability of an Attack

Probability	Equivalent
2^{-10}	full house in 5-card poker
2^{-20}	royal flush in 5-card poker
2^{-28}	you win this week's Powerball jackpot
2^{-40}	royal flush in 2 consecutive poker games
2^{-60}	the next meteorite that hits Earth lands on this slide

- While these examples are good to get some intuition, we need more concrete way to draw the line between *reasonable* and *unreasonable*.
- Just like running time, we will use an asymptotic approach to capture **low success probability** so that it can be tweaked as desired by changing the security parameter.

Negligible Functions

- A blind guessing adversary may “amplify” his success probability by guessing polynomial (say n^c) number of times, to achieve $\frac{n^c}{2^n}$ probability of success.

Negligible Functions

- A blind guessing adversary may “amplify” his success probability by guessing polynomial (say n^c) number of times, to achieve $\frac{n^c}{2^n}$ probability of success.
- This is still very low when n is large. Indeed, $\frac{1}{2^n}$ approaches zero so fast that it cannot be “rescued” by any polynomial

Negligible Functions

- A blind guessing adversary may “amplify” his success probability by guessing polynomial (say n^c) number of times, to achieve $\frac{n^c}{2^n}$ probability of success.
- This is still very low when n is large. Indeed, $\frac{1}{2^n}$ approaches zero so fast that it cannot be “rescued” by any polynomial
- We want to formalize this property so as to rule out the possibility that a PPT (i.e., probabilistic polynomial-time) adversary is able to amplify very low success probability to “non-trivial” success probability

Negligible Functions

- A blind guessing adversary may “amplify” his success probability by guessing polynomial (say n^c) number of times, to achieve $\frac{n^c}{2^n}$ probability of success.
- This is still very low when n is large. Indeed, $\frac{1}{2^n}$ approaches zero so fast that it cannot be “rescued” by any polynomial
- We want to formalize this property so as to rule out the possibility that a PPT (i.e., probabilistic polynomial-time) adversary is able to amplify very low success probability to “non-trivial” success probability

Negligible Functions

- A blind guessing adversary may “amplify” his success probability by guessing polynomial (say n^c) number of times, to achieve $\frac{n^c}{2^n}$ probability of success.
- This is still very low when n is large. Indeed, $\frac{1}{2^n}$ approaches zero so fast that it cannot be “rescued” by any polynomial
- We want to formalize this property so as to rule out the possibility that a PPT (i.e., probabilistic polynomial-time) adversary is able to amplify very low success probability to “non-trivial” success probability

Definition (Negligible Function)

A function $\nu(\cdot)$ is negligible if for every polynomial $p(\cdot)$, we have $\lim_{n \rightarrow \infty} p(n)\nu(n) = 0$

Negligible Functions (contd.)

- A negligible function approaches zero so fast that you can never catch up when multiplying by a polynomial

Negligible Functions (contd.)

- A negligible function approaches zero so fast that you can never catch up when multiplying by a polynomial
- **Alternatively:** A negligible function *decays faster than all “inverse-polynomial” functions*

Negligible Functions (contd.)

- A negligible function approaches zero so fast that you can never catch up when multiplying by a polynomial
- **Alternatively:** A negligible function *decays faster than all “inverse-polynomial” functions*

Negligible Functions (contd.)

- A negligible function approaches zero so fast that you can never catch up when multiplying by a polynomial
- **Alternatively:** A negligible function *decays faster than all “inverse-polynomial” functions*

Definition (Negligible Function)

A function $\nu(n)$ is negligible if $\forall c, \exists n_0$ such that $\forall n > n_0, \nu(n) \leq \frac{1}{n^c}$.

Negligible Functions: Examples

Problem

Let $f(\cdot)$ and $g(\cdot)$ be negligible functions. Show that $f(n) + g(n)$ is negligible.

Negligible Functions: Examples

Problem

Let $f(\cdot)$ and $g(\cdot)$ be negligible functions. Show that $f(n) + g(n)$ is negligible.

We need to show that $\forall c, \exists n_0$, such that $\forall n > n_0, f(n) + g(n) \leq \frac{1}{n^c}$.

Negligible Functions: Examples

Problem

Let $f(\cdot)$ and $g(\cdot)$ be negligible functions. Show that $f(n) + g(n)$ is negligible.

We need to show that $\forall c, \exists n_0$, such that $\forall n > n_0, f(n) + g(n) \leq \frac{1}{n^c}$.

- Since $f(\cdot)$ and $g(\cdot)$ are both negligible functions, we know that $\exists n_f, n_g$ corresponding to $c+1$, such that $\forall n > n_f, f(n) \leq \frac{1}{n^{c+1}}$ and $\forall n > n_g, g(n) \leq \frac{1}{n^{c+1}}$.

Negligible Functions: Examples

Problem

Let $f(\cdot)$ and $g(\cdot)$ be negligible functions. Show that $f(n) + g(n)$ is negligible.

We need to show that $\forall c, \exists n_0$, such that $\forall n > n_0, f(n) + g(n) \leq \frac{1}{n^c}$.

- Since $f(\cdot)$ and $g(\cdot)$ are both negligible functions, we know that $\exists n_f, n_g$ corresponding to $c+1$, such that $\forall n > n_f, f(n) \leq \frac{1}{n^{c+1}}$ and $\forall n > n_g, g(n) \leq \frac{1}{n^{c+1}}$.

For a given c , let $n_0 = \max(n_f, n_g, 2)$. $\forall n > n_0$:

$$\begin{aligned} f(n) + g(n) &\leq \frac{1}{n^{c+1}} + \frac{1}{n^{c+1}} \\ &\leq \frac{2}{n^{c+1}} \\ &\leq \frac{n}{n^{c+1}} \quad (\text{Since } n \geq n_0 \geq 2) \\ &\leq \frac{1}{n^c} \end{aligned}$$

Negligible Functions: Examples

Problem

Let $\nu(\cdot)$ be a negligible function and $p(\cdot)$ be a polynomial s.t. $p(n) \geq 0$, $\forall n > 0$. Show that $\nu(n) \cdot p(n)$ is negligible.

Negligible Functions: Examples

Problem

Let $\nu(\cdot)$ be a negligible function and $p(\cdot)$ be a polynomial s.t. $p(n) \geq 0$, $\forall n > 0$. Show that $\nu(n) \cdot p(n)$ is negligible.

We need to show that $\forall c$, $\exists n_0$, such that $\forall n > n_0$, $\nu(n) \cdot p(n) \leq \frac{1}{n^c}$.

Negligible Functions: Examples

Problem

Let $\nu(\cdot)$ be a negligible function and $p(\cdot)$ be a polynomial s.t. $p(n) \geq 0$, $\forall n > 0$. Show that $\nu(n) \cdot p(n)$ is negligible.

We need to show that $\forall c$, $\exists n_0$, such that $\forall n > n_0$, $\nu(n) \cdot p(n) \leq \frac{1}{n^c}$.

- Since $p(\cdot)$ is a polynomial function, we know that $\exists n_p, c_p$, such that, $\forall n > n_p$, $p(n) \leq n^{c_p}$.

Negligible Functions: Examples

Problem

Let $\nu(\cdot)$ be a negligible function and $p(\cdot)$ be a polynomial s.t. $p(n) \geq 0$, $\forall n > 0$. Show that $\nu(n) \cdot p(n)$ is negligible.

We need to show that $\forall c$, $\exists n_0$, such that $\forall n > n_0$, $\nu(n) \cdot p(n) \leq \frac{1}{n^c}$.

- Since $p(\cdot)$ is a polynomial function, we know that $\exists n_p, c_p$, such that, $\forall n > n_p$, $p(n) \leq n^{c_p}$.
- Since $\nu(\cdot)$ is a negligible function, we know that $\exists n_\nu$ corresponding to $c + c_p$, such that $\forall n > n_\nu$, $\nu(n) \leq \frac{1}{n^{c+c_p}}$.

Negligible Functions: Examples

Problem

Let $\nu(\cdot)$ be a negligible function and $p(\cdot)$ be a polynomial s.t. $p(n) \geq 0$, $\forall n > 0$. Show that $\nu(n) \cdot p(n)$ is negligible.

We need to show that $\forall c$, $\exists n_0$, such that $\forall n > n_0$, $\nu(n) \cdot p(n) \leq \frac{1}{n^c}$.

- Since $p(\cdot)$ is a polynomial function, we know that $\exists n_p, c_p$, such that, $\forall n > n_p$, $p(n) \leq n^{c_p}$.
- Since $\nu(\cdot)$ is a negligible function, we know that $\exists n_\nu$ corresponding to $c + c_p$, such that $\forall n > n_\nu$, $\nu(n) \leq \frac{1}{n^{c+c_p}}$.

For a given c , let $n_0 = \max(n_\nu, n_p)$. $\forall n > n_0$:

$$\begin{aligned}\nu(n) \cdot p(n) &\leq \frac{1}{n^{c+c_p}} \cdot n^{c_p} \\ &\leq \frac{1}{n^{c+c_p-c_p}} \\ &\leq \frac{1}{n^c}\end{aligned}$$

Distributions & Ensembles

- Recall: X is a distribution over sample space \mathcal{S} if it assigns probability p_s to the element $s \in \mathcal{S}$ s.t. $\sum_s p_s = 1$

Ensemble

A sequence $\{X_n\}_{n \in \mathbb{N}}$ is called an ensemble if for each $n \in \mathbb{N}$, X_n is a probability distribution over $\{0, 1\}^*$.

- Generally, X_n will be a distribution over the sample space $\{0, 1\}^{\ell(n)}$ (where $\ell(\cdot)$ is a polynomial)

Computational Indistinguishability

- So far in the definitions of uniform ciphertext security and perfect security, we required the two distributions to be **identical**.

Computational Indistinguishability

- So far in the definitions of uniform ciphertext security and perfect security, we required the two distributions to be **identical**.
- Going forward, since we (primarily) only care about security against efficient adversaries, our security definitions will not be quite as demanding.

Computational Indistinguishability

- So far in the definitions of uniform ciphertext security and perfect security, we required the two distributions to be **identical**.
- Going forward, since we (primarily) only care about security against efficient adversaries, our security definitions will not be quite as demanding.
- In most cases, it would suffice for two distributions (say X and Y) to “look alike” to *any* efficient test.

Computational Indistinguishability

- So far in the definitions of uniform ciphertext security and perfect security, we required the two distributions to be **identical**.
- Going forward, since we (primarily) only care about security against efficient adversaries, our security definitions will not be quite as demanding.
- In most cases, it would suffice for two distributions (say X and Y) to “look alike” to *any* efficient test.
- Efficient test = efficient computation = non-uniform PPT

Computational Indistinguishability

- So far in the definitions of uniform ciphertext security and perfect security, we required the two distributions to be **identical**.
- Going forward, since we (primarily) only care about security against efficient adversaries, our security definitions will not be quite as demanding.
- In most cases, it would suffice for two distributions (say X and Y) to “look alike” to *any* efficient test.
- Efficient test = efficient computation = non-uniform PPT
- If no non-uniform PPT adversary \mathcal{A} can tell them apart, we say that X and Y are **computationally indistinguishable**.

Computational Indistinguishability

- So far in the definitions of uniform ciphertext security and perfect security, we required the two distributions to be **identical**.
- Going forward, since we (primarily) only care about security against efficient adversaries, our security definitions will not be quite as demanding.
- In most cases, it would suffice for two distributions (say X and Y) to “look alike” to *any* efficient test.
- Efficient test = efficient computation = non-uniform PPT
- If no non-uniform PPT adversary \mathcal{A} can tell them apart, we say that X and Y are **computationally indistinguishable**.

Computational Indistinguishability

- So far in the definitions of uniform ciphertext security and perfect security, we required the two distributions to be **identical**.
- Going forward, since we (primarily) only care about security against efficient adversaries, our security definitions will not be quite as demanding.
- In most cases, it would suffice for two distributions (say X and Y) to “look alike” to *any* efficient test.
- Efficient test = efficient computation = non-uniform PPT
- If no non-uniform PPT adversary \mathcal{A} can tell them apart, we say that X and Y are **computationally indistinguishable**.

How do we formalize this?

Computational Indistinguishability



- \mathcal{A} 's output can be encoded using just one bit:
1 = “from X ” and 0 = “from Y ”

Computational Indistinguishability



- \mathcal{A} 's output can be encoded using just one bit:
1 = “from X ” and 0 = “from Y ”
- We want \mathcal{A} to output 1, with “almost similar” probability in both the above scenarios.

Computational Indistinguishability



- \mathcal{A} 's output can be encoded using just one bit:
1 = “from X ” and 0 = “from Y ”
- We want \mathcal{A} to output 1, with “almost similar” probability in both the above scenarios.

Computational Indistinguishability



- \mathcal{A} 's output can be encoded using just one bit:
1 = “from X ” and 0 = “from Y ”
- We want \mathcal{A} to output 1, with “almost similar” probability in both the above scenarios.

$$\Pr[x \leftarrow X; \mathcal{A}(1^n, x) = 1] \approx \Pr[y \leftarrow Y; \mathcal{A}(1^n, y) = 1] \implies$$

$$\left| \Pr[x \leftarrow X; \mathcal{A}(1^n, x) = 1] - \Pr[y \leftarrow Y; \mathcal{A}(1^n, y) = 1] \right| \leq \nu(n).$$

Computationally Indistinguishability: Definition

Definition (Computationally Indistinguishability)

Two ensembles of probability distributions $X = \{X_n\}_{n \in \mathbb{N}}$ and $Y = \{Y_n\}_{n \in \mathbb{N}}$ are said to be **computationally indistinguishable** if for every non-uniform PPT \mathcal{A} there exists a negligible function $\nu(\cdot)$ s.t.:

$$\left| \Pr[x \leftarrow X_n; \mathcal{A}(1^n, x) = 1] - \Pr[y \leftarrow Y_n; \mathcal{A}(1^n, y) = 1] \right| \leq \nu(n).$$

Computationally Indistinguishability: Definition

Definition (Computationally Indistinguishability)

Two ensembles of probability distributions $X = \{X_n\}_{n \in \mathbb{N}}$ and $Y = \{Y_n\}_{n \in \mathbb{N}}$ are said to be **computationally indistinguishable** if for every non-uniform PPT \mathcal{A} there exists a negligible function $\nu(\cdot)$ s.t.:

$$\left| \Pr[x \leftarrow X_n; \mathcal{A}(1^n, x) = 1] - \Pr[y \leftarrow Y_n; \mathcal{A}(1^n, y) = 1] \right| \leq \nu(n).$$

- The quantity $\left| \Pr[x \leftarrow X_n; D(1^n, x) = 1] - \Pr[y \leftarrow Y_n; D(1^n, y) = 1] \right|$ is called the **advantage** or bias of \mathcal{A} in distinguishing X and Y .

Computationally Indistinguishability: Definition

Definition (Computationally Indistinguishability)

Two ensembles of probability distributions $X = \{X_n\}_{n \in \mathbb{N}}$ and $Y = \{Y_n\}_{n \in \mathbb{N}}$ are said to be **computationally indistinguishable** if for every non-uniform PPT \mathcal{A} there exists a negligible function $\nu(\cdot)$ s.t.:

$$\left| \Pr[x \leftarrow X_n; \mathcal{A}(1^n, x) = 1] - \Pr[y \leftarrow Y_n; \mathcal{A}(1^n, y) = 1] \right| \leq \nu(n).$$

- The quantity $\left| \Pr[x \leftarrow X_n; D(1^n, x) = 1] - \Pr[y \leftarrow Y_n; D(1^n, y) = 1] \right|$ is called the **advantage** or bias of \mathcal{A} in distinguishing X and Y .
- Therefore, X and Y are computationally indistinguishable if all non-uniform PPT \mathcal{A} have negligible advantage in distinguishing them.

Properties of Computational Indistinguishability

- Notation: $\{X_n\} \approx_c \{Y_n\}$ means computational indistinguishability

Properties of Computational Indistinguishability

- Notation: $\{X_n\} \approx_c \{Y_n\}$ means computational indistinguishability
- **Closure:** If we apply an efficient operation on X and Y , they remain computationally indistinguishable. That is, \forall non-uniform PPT M

$$\{X_n\} \approx_c \{Y_n\} \implies \{M(X_n)\} \approx_c \{M(Y_n)\}$$

Properties of Computational Indistinguishability

- Notation: $\{X_n\} \approx_c \{Y_n\}$ means computational indistinguishability
- **Closure:** If we apply an efficient operation on X and Y , they remain computationally indistinguishable. That is, \forall non-uniform PPT M

$$\{X_n\} \approx_c \{Y_n\} \implies \{M(X_n)\} \approx_c \{M(Y_n)\}$$

Properties of Computational Indistinguishability

- **Notation:** $\{X_n\} \approx_c \{Y_n\}$ means computational indistinguishability
- **Closure:** If we apply an efficient operation on X and Y , they remain computationally indistinguishable. That is, \forall non-uniform PPT M

$$\{X_n\} \approx_c \{Y_n\} \implies \{M(X_n)\} \approx_c \{M(Y_n)\}$$

Proof Idea: If not, \mathcal{A} can use M to tell them apart!

Properties of Computational Indistinguishability

- **Notation:** $\{X_n\} \approx_c \{Y_n\}$ means computational indistinguishability
- **Closure:** If we apply an efficient operation on X and Y , they remain computationally indistinguishable. That is, \forall non-uniform PPT M

$$\{X_n\} \approx_c \{Y_n\} \implies \{M(X_n)\} \approx_c \{M(Y_n)\}$$

Proof Idea: If not, \mathcal{A} can use M to tell them apart!

- **Transitivity:** If X, Y are computationally indistinguishable, and Y, Z are computationally indistinguishable; then X, Z are also computationally indistinguishable.

Generalizing Transitivity: Hybrid Lemma

Lemma (Hybrid Lemma)

Let X^1, \dots, X^m be distribution ensembles for $m = \text{poly}(n)$. If for every $i \in [m - 1]$, X^i and X^{i+1} are computationally indistinguishable, then X^1 and X^m are computationally indistinguishable.

Generalizing Transitivity: Hybrid Lemma

Lemma (Hybrid Lemma)

Let X^1, \dots, X^m be distribution ensembles for $m = \text{poly}(n)$. If for every $i \in [m - 1]$, X^i and X^{i+1} are computationally indistinguishable, then X^1 and X^m are computationally indistinguishable.

This is the hybrid technique, stated more generally, in the computational setting.

Generalizing Transitivity: Hybrid Lemma


Lemma (Hybrid Lemma)

Let X^1, \dots, X^m be distribution ensembles for $m = \text{poly}(n)$. If for every $i \in [m - 1]$, X^i and X^{i+1} are computationally indistinguishable, then X^1 and X^m are computationally indistinguishable.

This is the hybrid technique, stated more generally, in the computational setting.

Used in most crypto proofs!

Hybrid Lemma



SmBC-comics.com

Looking Back at OTP

One-Time Pad

- $\text{KeyGen}(1^n) := k \xleftarrow{\$} \{0, 1\}^n$
- $\text{Enc}(k, m) := c = k \oplus m$
- $\text{Dec}(k, c) := m = k \oplus c$

Uniform ciphertext security and perfect security of OTP crucially depends on the following features:

- ① A key is as long as a message.
- ② A key is sampled uniformly at random.

Looking Back at OTP

One-Time Pad

- $\text{KeyGen}(1^n) := k \xleftarrow{\$} \{0, 1\}^n$
- $\text{Enc}(k, m) := c = k \oplus m$
- $\text{Dec}(k, c) := m = k \oplus c$

Uniform ciphertext security and perfect security of OTP crucially depends on the following features:

- ① A key is as long as a message.
- ② A key is sampled uniformly at random.

These features of OTP are also its key drawbacks! Why?

Drawbacks of OTP

- **Length of the key is as much as that of the message.**
 - The key somehow needs to be exchanged between the sender and recipient in advance.
 - Longer the key, more inefficient the key exchange process will be.

Drawbacks of OTP

- **Length of the key is as much as that of the message.**
 - The key somehow needs to be exchanged between the sender and recipient in advance.
 - Longer the key, more inefficient the key exchange process will be.
- **The key is sampled uniformly at random.**
 - Common sources of randomness:
 - key-strokes
 - mouse movement
 - power consumption
 - ...
 - These processes can only produce so much true randomness

Drawbacks of OTP

- **Length of the key is as much as that of the message.**
 - The key somehow needs to be exchanged between the sender and recipient in advance.
 - Longer the key, more inefficient the key exchange process will be.
- **The key is sampled uniformly at random.**
 - Common sources of randomness:
 - key-strokes
 - mouse movement
 - power consumption
 - ...
 - These processes can only produce so much true randomness

Can we **expand** few random bits into many random “looking” bits?

Pseudorandomness

- Suppose you have n uniformly random bits: $x = x_1 \| \dots \| x_n$

Pseudorandomness

- Suppose you have n uniformly random bits: $x = x_1 \| \dots \| x_n$
- Find a **deterministic** (polynomial-time) algorithm G such that:

Pseudorandomness

- Suppose you have n uniformly random bits: $x = x_1 \| \dots \| x_n$
- Find a **deterministic** (polynomial-time) algorithm G such that:
 - $G(x)$ outputs (say) $n + 1$ bits: $y = y_1 \| \dots \| y_{n+1}$

Pseudorandomness

- Suppose you have n uniformly random bits: $x = x_1 \| \dots \| x_n$
- Find a **deterministic** (polynomial-time) algorithm G such that:
 - $G(x)$ outputs (say) $n + 1$ bits: $y = y_1 \| \dots \| y_{n+1}$
 - y looks “as good as” a truly random string $r = r_1 \| \dots \| r_{n+1}$

Pseudorandomness

- Suppose you have n uniformly random bits: $x = x_1 \| \dots \| x_n$
- Find a **deterministic** (polynomial-time) algorithm G such that:
 - $G(x)$ outputs (say) $n + 1$ bits: $y = y_1 \| \dots \| y_{n+1}$
 - y looks “as good as” a truly random string $r = r_1 \| \dots \| r_{n+1}$
- That is, the following distributions are computationally indistinguishable.

$$\left\{ x \xleftarrow{\$} \{0, 1\}^n : G(x) \right\} \quad \text{and} \quad \left\{ r \xleftarrow{\$} \{0, 1\}^{n+1} : r \right\}$$

Pseudorandomness

- Suppose you have n uniformly random bits: $x = x_1 \| \dots \| x_n$
- Find a **deterministic** (polynomial-time) algorithm G such that:
 - $G(x)$ outputs (say) $n + 1$ bits: $y = y_1 \| \dots \| y_{n+1}$
 - y looks “as good as” a truly random string $r = r_1 \| \dots \| r_{n+1}$
- That is, the following distributions are computationally indistinguishable.

$$\left\{ x \xleftarrow{\$} \{0, 1\}^n : G(x) \right\} \quad \text{and} \quad \left\{ r \xleftarrow{\$} \{0, 1\}^{n+1} : r \right\}$$

- $G : \{0, 1\}^n \rightarrow \{0, 1\}^{n+1}$ is called a **pseudorandom generator** (PRG) and its output is a **pseudorandom** string.

Pseudorandom Generators (PRG)


Definition (Pseudorandom Generator)

A deterministic algorithm G is called a pseudorandom generator (PRG) if:

- G can be computed in polynomial time
- $|G(x)| > |x|$
- $\left\{ x \xleftarrow{\$} \{0, 1\}^n; G(x) \right\} \approx_c \left\{ U_{\ell(n)} \right\}$ where $\ell(n) = |G(0^n)|$


The **stretch** of G is defined as: $|G(x)| - |x|$

Illustrating PRG $G : \{0, 1\}^\lambda \rightarrow \{0, 1\}^{2\lambda}$



- From a relative perspective, the PRG's output distribution is tiny. Out of the $2^{2\lambda}$ strings in $\{0, 1\}^{2\lambda}$, only 2^λ are possible outputs of G . These strings make up $2^\lambda/2^{2\lambda} = 1/2^\lambda$ fraction of $\{0, 1\}^{2\lambda}$ — a **negligible fraction!**

Illustrating PRG $G : \{0, 1\}^\lambda \rightarrow \{0, 1\}^{2\lambda}$



- From a relative perspective, the PRG's output distribution is tiny. Out of the $2^{2\lambda}$ strings in $\{0, 1\}^{2\lambda}$, only 2^λ are possible outputs of G . These strings make up $2^\lambda/2^{2\lambda} = 1/2^\lambda$ fraction of $\{0, 1\}^{2\lambda}$ — a **negligible fraction!**
- From an absolute perspective, the PRG's output distribution is huge. There are 2^λ possible outputs of G , which is an **exponential amount!** This is large enough that an efficient adversary cannot distinguish it from the set $\{0, 1\}^{2\lambda}$.

Pseudorandom OTP

- We can now use a PRG to modify one-time pads as follows.

Pseudorandom One-Time Pad

Let n be the security parameter and $\ell(\cdot)$ be a polynomial. Let $G : \{0, 1\}^n \rightarrow \{0, 1\}^{\ell(n)}$ be a PRG, and let the message space and ciphertext space be $\{0, 1\}^{\ell(n)}$.

- $\text{KeyGen}(1^n) := k \xleftarrow{\$} \{0, 1\}^n$
- $\text{Enc}(k, m) := c = G(k) \oplus m$
- $\text{Dec}(k, c) := m = G(k) \oplus c$

Pseudorandom OTP

- We can now use a PRG to modify one-time pads as follows.

Pseudorandom One-Time Pad

Let n be the security parameter and $\ell(\cdot)$ be a polynomial. Let $G : \{0, 1\}^n \rightarrow \{0, 1\}^{\ell(n)}$ be a PRG, and let the message space and ciphertext space be $\{0, 1\}^{\ell(n)}$.

- $\text{KeyGen}(1^n) := k \xleftarrow{\$} \{0, 1\}^n$
- $\text{Enc}(k, m) := c = G(k) \oplus m$
- $\text{Dec}(k, c) := m = G(k) \oplus c$

Does this encryption scheme satisfy our original definitions of one-time uniform ciphertext security or one-time perfect security?